

# Algorithms and Data Structures for DS

## Encoding Number encoding

Input: 15

2	15		
2	7	—	1
2	3	—	1
	1	—	1

### Learning goals

- Codes for numbers

Binary number: 1111

# NUMBERS ON DIGITAL MACHINES

- Why do we care?
- Pretty much whatever we do in DS is implemented as some form of number crunching on a machine
- Computers perform billions of arithmetic operations per second
- But: Not all results are exact – even simple arithmetic may involve rounding or approximation
- Understanding how numbers are represented and processed is crucial for correct interpretation of results

# NUMBERS ON DIGITAL MACHINES

- The basic arithmetic operations are performed directly by the CPU
- The fewer bits per number, the faster
- Performance tradeoff: more precision often means slower computation
- To be efficient, numbers are encoded with a **fixed number of bytes**, so with  $N$  bits
- So we map infinite sets  $\mathbb{Z}$  and  $\mathbb{R}$  to  $2^N$  machine numbers
- Such a mapping introduces rounding, overflow, and underflow
- Leads to potential loss of information.

# “BUG”-REPORT IN R I

To: R-bugs@biostat.ku.dk

Subject: error in trunc function

the command get a wrong result

```
> trunc(2.3 * 100)
```

```
[1] 229
```

Answer Duncan Murdoch:

That is the correct answer. 2.3 is not representable exactly; the actual value used is slightly less.

# “BUG”-REPORT IN R II

To: R-bugs@biostat.ku.dk

Subject: [Rd] match() (PR#13135)

The match function does not return value properly.

See an example below.

```
> a = seq(0.6, 1, by = 0.01)
```

```
> match(0.88, a)
```

```
[1] 29
```

```
> match(0.89, a)
```

```
[1] NA
```

Answer Brian Ripley:

FAQ Q7.31 strikes again!

0.89 is not a member of `seq(0.6,1,by=0.01)`, since 0.01 cannot be represented exactly in a binary computer.

# “BUG”-REPORT IN R III

Subject: Re: Bug in R?

> Hi, I'm not sure if it's really a bug:

> When you execute:

>> (2 / 3) \* (0.6 / (1 - 0.6))

> the result will be:

> [1] 1

> but if you execute:

>> (2 / 3) \* (0.6 / (1 - 0.6)) == 1

> the result is:

> [1] FALSE

> Note: I'm using version 2.9.2, (and tried it in 2.9.1 in 2.9.1 too)  
with Microsoft Windows XP [Version 5.1.2600].

FAQ 7.31 strikes again:

R> 1 - (2 / 3) \* (0.6 / (1 - 0.6))

[1] 2.220446e-16

R> .Machine\$double.eps

[1] 2.220446e-16

# “BUG”-REPORT IN R IV

```
> Dear all,  
>  
> might seem an easy question but I cannot figure it out.  
>  
> floor(100 * (.58))  
> [1] 57  
>  
> where is the trick here?  
> And how can I end up with the right answer?
```

See \texttt{R} FAQ 7.31

```
> sprintf("%.20f", 100 * .58)  
[1] "57.99999999999999289457"
```